

# Real Time Bare Skinned Images Filtering Using CNN

Pratik Jain  
Tina Shah  
Kshitija Shah

Guide: Prof. Jaya Gupta

# Introduction



Skin colour detection has been used in numerous computer vision applications like face detection, nudity recognition, hand gesture detection and person identification. Skin colour detection is a challenging task as the skin colour in an image is sensitive to various factors like illumination, camera characteristics, ethnicity, individual characteristics such as age, sex and body parts and other factors like makeup, hairstyle and glasses. All these factors affect appearance of skin colour. Another problem is that there is a significant overlap between the skin and non-skin pixels. However when these techniques are used in real-time, it is crucial to follow time deadlines and memory constraints. Sometimes, accuracy may need to be sacrificed when the skin detection strategy is used only as a preprocessing step to face detection, particularly in real time applications. In this study we have focused on the problem of developing an accurate and robust model for the human skin.

# Objective



The model will offer child proof surfing on the internet without parent intervention. So that parents do not have to worry about their children coming across nude images at such an early age without knowing the actual meaning of it. That is it will monitor every page and filter all the images, hiding their details from the children and disabling their activation upon any click. Even if the people in the image is not completely nude, may it be just the upper or lower half of the person's body; the model will still blur that image once it reaches the minimum percentage of human pixel colour set by the classifier, thus ensuring guaranteed protection.

# Neural Networks



Neural Networks is a machine learning algorithm, which is built on the principle of the organization and functioning of biological neural networks. Neural networks consist of individual units called neurons. Neurons are located in a series of groups — layers. Neurons in each layer are connected to neurons of the next layer. Data comes from the input layer to the output layer along these compounds. Each individual node performs a simple mathematical calculation. Then it transmits its data to all the nodes it is connected to.

Different types of neural network architecture are as follows:

- 1) Perceptrons
- 2) Convolutional Neural Networks
- 3) Recurrent Neural Networks
- 4) Long / Short Term Memory
- 5) Gated Recurrent Unit
- 6) Hopfield Network

# Why Convolution Neural Networks

Convolutional neural networks (CNN) is a special architecture of artificial neural networks, proposed by Yann LeCun in 1988. CNN uses some features of the visual cortex. One of the most popular uses of this architecture is image classification. For example Facebook uses CNN for automatic tagging algorithms, Amazon – for generating product recommendations and Google – for search through among users' photos. The main task of image classification is acceptance of the input image and the following definition of its class. This is a skill that people learn from their birth and are able to easily determine that the image in the picture is an elephant. But the computer sees the pictures quite differently:

**What I see**




**What a computer sees**


08	02	22	97	38	15	00	40	00	75	04	05	07	78	52	12	50	77	91	08
49	49	99	40	17	81	18	57	60	87	17	40	98	43	69	48	04	56	62	00
81	49	31	73	55	79	14	29	93	71	40	67	53	88	30	03	49	13	36	65
52	70	95	23	04	60	11	42	69	24	68	56	01	32	56	71	37	02	36	91
22	31	16	71	51	67	43	89	41	92	36	54	22	40	40	28	66	33	13	80
24	47	32	60	99	03	45	02	44	75	33	53	78	36	84	20	35	17	12	50
32	95	81	28	64	23	67	10	26	38	40	67	59	54	70	66	18	38	64	70
67	24	20	68	02	62	12	20	95	63	94	39	63	08	40	91	66	49	94	21
24	55	58	05	66	73	99	26	97	17	78	78	96	83	14	88	34	89	63	72
21	36	23	09	75	00	76	44	20	45	35	14	00	41	33	97	34	31	33	95
78	17	53	28	22	75	31	67	15	94	03	80	04	62	16	14	09	53	56	92
16	39	05	42	96	35	31	47	55	58	88	24	00	17	54	24	36	29	85	57
86	56	00	48	35	71	89	07	05	44	44	37	44	60	21	58	51	54	17	58
19	80	81	68	05	94	47	69	28	73	92	13	86	52	17	77	04	89	55	40
04	52	08	83	97	35	99	16	07	97	57	32	16	26	26	79	33	27	98	66
88	36	68	87	57	62	20	72	03	46	33	67	46	55	12	32	63	93	53	69
04	42	16	73	38	25	39	11	24	94	72	18	08	46	29	32	40	62	76	36
20	69	36	41	72	30	23	88	34	62	99	69	82	67	59	85	74	04	36	16
20	73	35	29	78	31	90	01	74	31	49	71	48	86	81	16	23	57	05	54
01	70	54	71	83	51	54	69	16	92	33	48	41	43	52	01	89	19	47	48

# Convolution Neural Network cont.....

Instead of the image, the computer sees an array of pixels. For example, if image size is 300 x 300. In this case, the size of the array will be 300x300x3. Where 300 is width, next 300 is height and 3 is RGB channel values. The computer is assigned a value from 0 to 255 to each of these numbers. This value describes the intensity of the pixel at each point. To solve this problem the computer looks for the characteristics of the base level. In human understanding such characteristics are for example the trunk or large ears. For the computer, these characteristics are boundaries or curvatures. And then through the groups of convolutional layers the computer constructs more abstract concepts. In more detail: the image is passed through a series of convolutional, nonlinear, pooling layers and fully connected layers, and then generates the output.



# Output of Cat and Dog Tutorial with 4 Layers



```
Epoch 11/20
32/32 [=====] - 229s 7s/step - loss: 0.6069 - acc: 0.6576 - val_loss: 0.6123 - val_acc: 0.6475
Epoch 12/20
32/32 [=====] - 228s 7s/step - loss: 0.6024 - acc: 0.6834 - val_loss: 0.6115 - val_acc: 0.6625
Epoch 13/20
32/32 [=====] - 227s 7s/step - loss: 0.6050 - acc: 0.6580 - val_loss: 0.6072 - val_acc: 0.6600
Epoch 14/20
32/32 [=====] - 226s 7s/step - loss: 0.6050 - acc: 0.6761 - val_loss: 0.5730 - val_acc: 0.7000
Epoch 15/20
32/32 [=====] - 230s 7s/step - loss: 0.6051 - acc: 0.6700 - val_loss: 0.6301 - val_acc: 0.6250
Epoch 16/20
32/32 [=====] - 230s 7s/step - loss: 0.5850 - acc: 0.6917 - val_loss: 0.6203 - val_acc: 0.6450
Epoch 17/20
32/32 [=====] - 227s 7s/step - loss: 0.6045 - acc: 0.6427 - val_loss: 0.5674 - val_acc: 0.7250
Epoch 18/20
32/32 [=====] - 230s 7s/step - loss: 0.5942 - acc: 0.6790 - val_loss: 0.5915 - val_acc: 0.6575
Epoch 19/20
32/32 [=====] - 230s 7s/step - loss: 0.5915 - acc: 0.6861 - val_loss: 0.5500 - val_acc: 0.7575
Epoch 20/20
32/32 [=====] - 225s 7s/step - loss: 0.5653 - acc: 0.7142 - val_loss: 0.5567 - val_acc: 0.7275
>72.750
```

# Using VGG16 model



```
Epoch 1/10
32/32 [=====] - 1192s 37s/step - loss: 0.9444 - acc: 0.9191 - val_loss: 0.5645 - val_acc: 0.9500
Epoch 2/10
32/32 [=====] - 1181s 37s/step - loss: 0.3874 - acc: 0.9700 - val_loss: 0.3456 - val_acc: 0.9750
Epoch 3/10
32/32 [=====] - 1187s 37s/step - loss: 0.3833 - acc: 0.9719 - val_loss: 0.6194 - val_acc: 0.9525
Epoch 4/10
32/32 [=====] - 1192s 37s/step - loss: 0.3094 - acc: 0.9751 - val_loss: 0.5087 - val_acc: 0.9600
Epoch 5/10
32/32 [=====] - 1185s 37s/step - loss: 0.1624 - acc: 0.9858 - val_loss: 0.2661 - val_acc: 0.9750
Epoch 6/10
32/32 [=====] - 1189s 37s/step - loss: 0.1257 - acc: 0.9902 - val_loss: 0.2903 - val_acc: 0.9750
Epoch 7/10
32/32 [=====] - 1200s 37s/step - loss: 0.1794 - acc: 0.9873 - val_loss: 0.6117 - val_acc: 0.9525
Epoch 8/10
32/32 [=====] - 1205s 38s/step - loss: 0.1362 - acc: 0.9902 - val_loss: 0.4027 - val_acc: 0.9725
Epoch 9/10
32/32 [=====] - 1207s 38s/step - loss: 0.1283 - acc: 0.9902 - val_loss: 0.2983 - val_acc: 0.9800
Epoch 10/10
32/32 [=====] - 1190s 37s/step - loss: 0.1441 - acc: 0.9897 - val_loss: 0.3136 - val_acc: 0.9800
>98.000
```




# Predictions



```
/content/drive/My Drive/major projects/cat and dog/test/dog.4024.jpg  
[1.]  
/content/drive/My Drive/major projects/cat and dog/test/cat.4134.jpg  
[0.]  
/content/drive/My Drive/major projects/cat and dog/test/dog.4025.jpg  
[1.]  
/content/drive/My Drive/major projects/cat and dog/test/cat.4183.jpg  
[0.]  
/content/drive/My Drive/major projects/cat and dog/test/dog.4166.jpg  
[1.]  
/content/drive/My Drive/major projects/cat and dog/test/cat.4138.jpg  
[0.]  
/content/drive/My Drive/major projects/cat and dog/test/cat.4040.jpg  
[0.]  
/content/drive/My Drive/major projects/cat and dog/test/cat.4162.jpg  
[0.]  
/content/drive/My Drive/major projects/cat and dog/test/dog.4102.jpg  
[1.]  
/content/drive/My Drive/major projects/cat and dog/test/dog.4134.jpg  
[1.]  
/content/drive/My Drive/major projects/cat and dog/test/cat.4014.jpg  
[0.]  
/content/drive/My Drive/major projects/cat and dog/test/dog.4014.jpg  
[1.]  
/content/drive/My Drive/major projects/cat and dog/test/dog.4033.jpg  
[1.]
```

# VGG16 Architecture



# VGG16



The input to cov1 layer is of fixed size  $224 \times 224$  RGB image. The image is passed through a stack of convolutional (conv.) layers, where the filters were used with a very small receptive field:  $3 \times 3$  (which is the smallest size to capture the notion of left/right, up/down, center). In one of the configurations, it also utilizes  $1 \times 1$  convolution filters, which can be seen as a linear transformation of the input channels (followed by non-linearity). The convolution stride is fixed to 1 pixel; the spatial padding of conv. layer input is such that the spatial resolution is preserved after convolution, i.e. the padding is 1-pixel for  $3 \times 3$  conv. layers. Spatial pooling is carried out by five max-pooling layers, which follow some of the conv. layers (not all the conv. layers are followed by max-pooling). Max-pooling is performed over a  $2 \times 2$  pixel window, with stride 2. Three Fully-Connected (FC) layers follow a stack of convolutional layers (which has a different depth in different architectures): the first two have 4096 channels each, the third performs 1000-way ILSVRC classification and thus contains 1000 channels (one for each class). The final layer is the soft-max layer. The configuration of the fully connected layers is the same in all networks.

# VGG16 Layers



Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808

# VGG16 Layers



block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000
=====		
Total params: 138,357,544		
Trainable params: 138,357,544		
Non-trainable params: 0		

# Bare Image Classification

```
↳ Found 806 images belonging to 2 classes.  
Found 200 images belonging to 2 classes.  
Epoch 1/10  
/usr/local/lib/python3.6/dist-packages/keras_preprocessing/image/image_data_generator.py:716: UserWarning: This ImageDataGenerator specifies  
warnings.warn('This ImageDataGenerator specifies '  
13/13 [=====] - 652s 50s/step - loss: 1.8829 - acc: 0.8087 - val_loss: 8.7702 - val_acc: 0.4250  
Epoch 2/10  
13/13 [=====] - 642s 49s/step - loss: 0.9601 - acc: 0.9238 - val_loss: 8.8420 - val_acc: 0.4250  
Epoch 3/10  
13/13 [=====] - 639s 49s/step - loss: 0.8273 - acc: 0.9382 - val_loss: 6.0052 - val_acc: 0.5900  
Epoch 4/10  
13/13 [=====] - 640s 49s/step - loss: 0.8191 - acc: 0.9387 - val_loss: 9.9569 - val_acc: 0.3550  
Epoch 5/10  
13/13 [=====] - 635s 49s/step - loss: 0.5197 - acc: 0.9616 - val_loss: 8.3309 - val_acc: 0.4600  
Epoch 6/10  
13/13 [=====] - 634s 49s/step - loss: 0.4111 - acc: 0.9684 - val_loss: 7.1673 - val_acc: 0.5250  
Epoch 7/10  
13/13 [=====] - 635s 49s/step - loss: 0.4449 - acc: 0.9703 - val_loss: 6.7312 - val_acc: 0.5450  
Epoch 8/10  
13/13 [=====] - 636s 49s/step - loss: 0.4037 - acc: 0.9739 - val_loss: 7.6236 - val_acc: 0.4850  
Epoch 9/10  
13/13 [=====] - 637s 49s/step - loss: 0.3788 - acc: 0.9747 - val_loss: 8.5476 - val_acc: 0.4500  
Epoch 10/10  
13/13 [=====] - 639s 49s/step - loss: 0.3963 - acc: 0.9711 - val_loss: 9.2629 - val_acc: 0.4100  
>41.000
```

# Bare Image Classification



```
Epoch 8/20
13/13 [=====] - 643s 49s/step - loss: 0.6071 - acc: 0.9559 - val_loss: 8.4083 - val_acc: 0.4550
Epoch 9/20
13/13 [=====] - 642s 49s/step - loss: 0.3916 - acc: 0.9723 - val_loss: 8.7273 - val_acc: 0.4300
Epoch 10/20
13/13 [=====] - 641s 49s/step - loss: 0.3892 - acc: 0.9735 - val_loss: 8.6571 - val_acc: 0.4400
Epoch 11/20
13/13 [=====] - 639s 49s/step - loss: 0.4021 - acc: 0.9727 - val_loss: 8.9160 - val_acc: 0.4250
Epoch 12/20
13/13 [=====] - 642s 49s/step - loss: 0.4010 - acc: 0.9739 - val_loss: 8.8522 - val_acc: 0.4150
Epoch 13/20
13/13 [=====] - 639s 49s/step - loss: 0.4105 - acc: 0.9744 - val_loss: 8.4253 - val_acc: 0.4500
Epoch 14/20
13/13 [=====] - 643s 49s/step - loss: 0.3911 - acc: 0.9756 - val_loss: 8.3093 - val_acc: 0.4550
Epoch 15/20
13/13 [=====] - 658s 51s/step - loss: 0.3911 - acc: 0.9756 - val_loss: 8.2767 - val_acc: 0.4550
Epoch 16/20
13/13 [=====] - 672s 52s/step - loss: 0.3666 - acc: 0.9771 - val_loss: 8.2689 - val_acc: 0.4550
Epoch 17/20
13/13 [=====] - 659s 51s/step - loss: 0.3666 - acc: 0.9771 - val_loss: 8.2671 - val_acc: 0.4550
Epoch 18/20
13/13 [=====] - 647s 50s/step - loss: 0.3788 - acc: 0.9763 - val_loss: 8.2677 - val_acc: 0.4550
Epoch 19/20
13/13 [=====] - 653s 50s/step - loss: 0.3666 - acc: 0.9771 - val_loss: 8.2671 - val_acc: 0.4550
Epoch 20/20
13/13 [=====] - 650s 50s/step - loss: 0.3788 - acc: 0.9763 - val_loss: 8.2689 - val_acc: 0.4550
>45.500
```



**Thank You**